PRINCE GOUR

My Portfolio | gourprince2004@gmail.com | github | linkedin

PROFILE

Results-oriented AI & Generative AI Engineer with a strong foundation in developing and fine-tuning Large Language Models (LLMs), transformer-based architectures, and reinforcement learning agents. Skilled in applying deep learning solutions to real-world NLP and computer vision tasks.

EDUCATION

Poornima University – Jaipur, India

Bachelor of Computer Applications 2022-2025

SKILLS

Technical Skills: Python, PyTorch, TensorFlow, Transformers, Deep Learning, Natural Language Processing (NLP), Reinforcement Learning, Retrieval-Augmented Generation (RAG), LLM Fine-tuning (LoRA, PEFT), Prompt Engineering.

Tools & Frameworks: LangChain, LangGraph, AI Agentic Frameworks (Autogen, CrewAI), FastAPI, Docker, Git, Weights & Biases (W&B), Hugging Face Transformers.

Languages & Scripting: Python, TypeScript, SQL, Bash, C++, C#.

EXPERIENCE

NeuroNex Labs | Jaipur, India — Jan 2025 – Nov 2025

AI Engineer

- Fine-tuned domain-specificLLMs using LoRA/PEFT techniques to enable low-latency inference.
- Managed and optimized end-to-end model pipelines including dataset curation, evaluation, and benchmarking.

Shreshtha Consulatant | Jaipur, India — Dec 2025 – present

AI Engineer

- Developed and trained custom AI models to automate company-specific MEP tasks, reducing manual effort and improving workflow efficiency.
- Built AI-powered tools and intelligent ZWCAD plugins to streamline designer workflows and enhance productivity across engineering design processes.

PROJECTS

- **Ternary Neural Network (MNIST):** Developed a ternary-weight neural network for MNIST digit classification using weight quantization (–1,0,+1) for efficient and lightweight inference.
- **RLHF Agent:** Developed a reinforcement learning agent fine-tuned using preference-based human feedback and PPO.
- **Stable Diffusion Pipeline:** Rebuilt Stable Diffusion architecture and integrated official weights for high-quality image generation.
- **MOE [Mistral 8x7B]:** Implemented a Mixture-of-Experts architecture inspired by Mistral 8x7B, optimizing expert routing and load balancing for efficient inference.